

Machine Learning Assginment -4

Name : Sachit Kumar Tadishetty
700# : 700734682
UCM ID : SXT46820

VideoLink:

https://drive.google.com/drive/folders/1xq28VUjgRu-2ng9uIL5182B9izNjOkDe?usp=share_link

Github Link : <https://github.com/sachit46820/ML-Assignment>

1) Apply Linear Regression to the provided dataset using underlying steps.

- Import the given “Salary_Data.csv”
- Split the data in train_test partitions, such that 1/3 of the data is reserved as test subset.
- Train and predict the model.
- Calculate the mean_squared error
- Visualize both train and test data using scatter plot.

STUDENT / ML Assignment 4 (SACHIT) / Notebook 1

● Ready

▶ Run notebook ▼ ↺

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
from sklearn.decomposition import PCA
from sklearn.preprocessing import LabelEncoder, StandardScaler
import seaborn as sns
sns.set(style="white", color_codes=True)
import warnings
warnings.filterwarnings("ignore")
```

[1]

- Used Linear regression to find the relationship between variables.
- Imported simpleimputer which helps in handling the missing data in the predictive dataset.
- Imported Kmeans for clustering the data points into K-clusters by minimising the variance.
- Imported metrics for evaluating the algorithms.

warnings.filterwarnings('ignore')

Ready

Run notebook

⋮

```
df=pd.read_csv("Salary_Data.csv")
df.head()
```

[2]

	YearsExperience f. ▾	Salary float64 ▾	
0	1.1	39343.0	
1	1.3	46205.0	
2	1.5	37731.0	
3	2.0	43525.0	
4	2.2	39891.0	

Visualize

5 rows, showing 10 per page

<< < Page 1 of 1 >

↓

```
X = df.iloc[:, :-1].values
Y = df.iloc[:, 1].values
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X,Y, test_size=1/3,random_state = 0)
```

[3]

```
regressor = LinearRegression()
regressor.fit(X_Train, Y_Train)

Y_Pred = regressor.predict(X_Test)
```

[4]

Y_Pred = regressor.predict(X_Test)

Ready

Run notebook

mean_squared_error(Y_Test,Y_Pred)

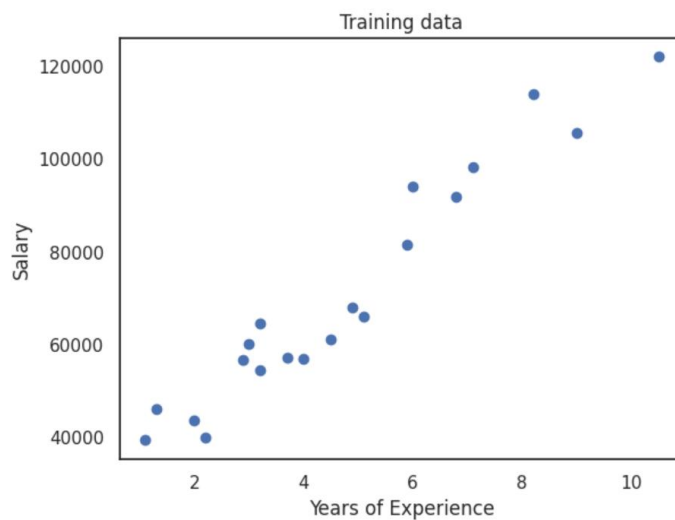
[5]

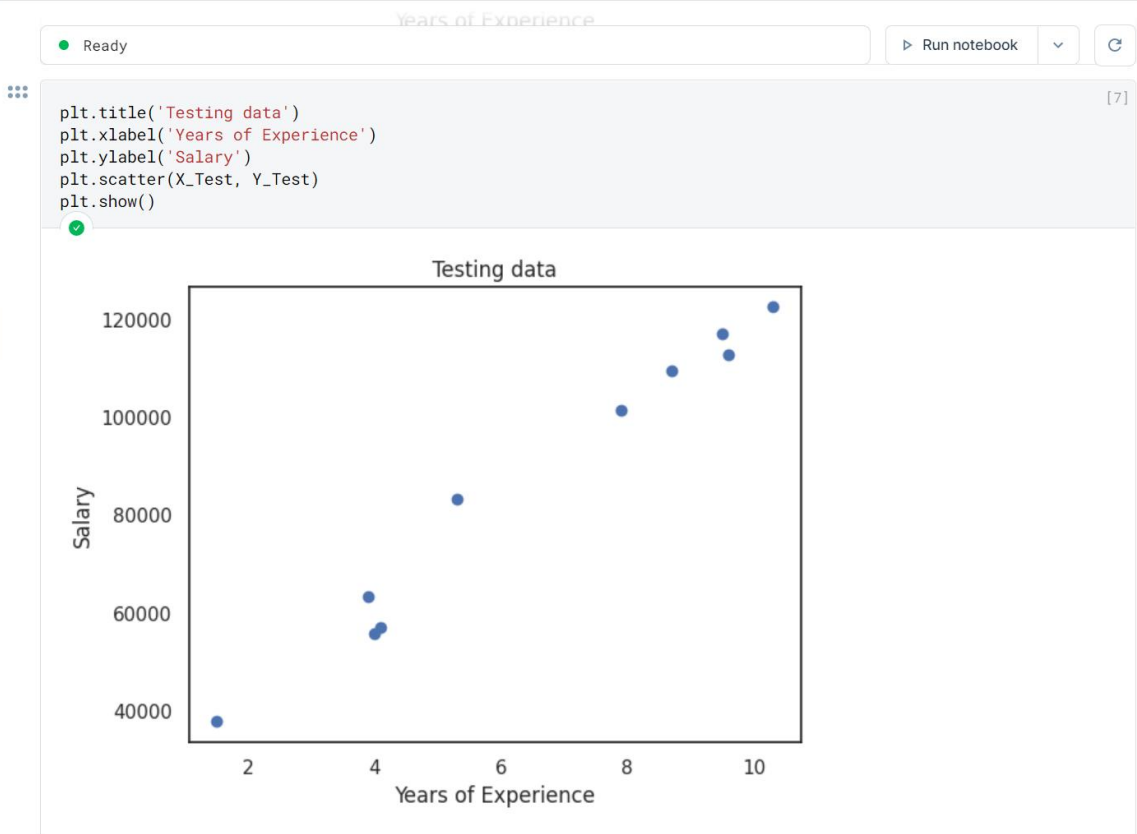
21026037.329511296

⋮

```
plt.title('Training data')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.scatter(X_Train, Y_Train)
plt.show()
```

[6]





2) Apply K means clustering in the dataset provided:

- Remove any null values by the mean.
- Use the elbow method to find a good number of clusters with the K-Means algorithm
- Calculate the silhouette score for the above clustering.

Ready

Run notebook



```
df2=pd.read_csv("K-Mean_Dataset.csv")
df2.head()
```

[8]

	CUST_ID object	BALANCE float64	BALANCE_FREQ...	PURCHASES floa...	ONEOFF_PURC...	INSTALLMENTS...	CASH_AI
0	C10001	40.900749	0.818182	95.4	0.0	95.4	
1	C10002	3202.467416	0.909091	0.0	0.0	0.0	64
2	C10003	2495.148862	1.0	773.17	773.17	0.0	
3	C10004	1666.670542	0.636364	1499.0	1499.0	0.0	2
4	C10005	817.714335	1.0	16.0	16.0	0.0	

```
X = df2.iloc[:,1:].values

imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
imputer = imputer.fit(X)
X = imputer.transform(X)
```

[9]

X = imputer.transform(X)

Ready


Run notebook



```
wcss = []
for i in range(1,11):
    kmeans = KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=10,random_state=0)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

plt.plot(range(1,11),wcss)
plt.title('the elbow method')
plt.xlabel('Number of Clusters')
plt.ylabel('Wcss')
plt.show()
```

[10]



STUDENT / ML Assignment 4 (SACHIT) / Notebook 1

Number of Clusters

Ready

Run notebook

```
1 from sklearn.cluster import KMeans
2 nclusters = 4 # this is the k in kmeans
3 km = KMeans(n_clusters=nclusters)
4 km.fit(X)
```

KMeans

KMeans(n_clusters=4)

```
y_cluster_kmeans = km.predict(X)
from sklearn import metrics
score = metrics.silhouette_score(X, y_cluster_kmeans)
print('Silhouette score:', score)
```

Silhouette score: 0.4657118789980141

3) Try feature scaling and then apply K-Means on the scaled features. Did that improve the Silhouette score? If Yes, can you justify why

STUDENT / ML Assignment 4 (SACHIT) / Notebook 1

Ready

Run notebook

```
scaler = preprocessing.StandardScaler()
scaler.fit(X)
X_scaled_array = scaler.transform(X)
X_scaled = pd.DataFrame(X_scaled_array,)
```

```
from sklearn.cluster import KMeans
nclusters = 4
km = KMeans(n_clusters=nclusters)
km.fit(X_scaled)
```

KMeans

KMeans(n_clusters=4)

```
y_scaled_cluster_kmeans = km.predict(X_scaled)
from sklearn import metrics
score = metrics.silhouette_score(X_scaled, y_scaled_cluster_kmeans)
print('Silhouette score after applying scaling:', score)
```

Silhouette score after applying scaling: 0.1976074492720698

Code Text SQL Chart Input

No, the Silhouette score did not improve after feature scaling. because scaling will basically result in models that are disproportionally influenced by the subset of features on a large scale hence the performance decreases.